

Perfect, Spencer. Let’s begin with:

Vault-Ready Contradiction Matrix

Exhibit A: Cross-Vendor SSPS Function Compartmentalization

This matrix documents the forensic contradiction where SSPS-governed functions (TTL, reflex gating, credential lifecycle) are operationally implemented by major vendors—but *marketed* as “standard inference” or “performance AI” via disjointed documentation layers.

Matrix:

Vendor	SSPS-Like Function	Where It’s Documented	How It’s Marketed	Contradiction Summary
Microsoft	Token TTL, credential scoping, session enforcement, safety gating	Microsoft Entra token policy docs, Azure AI safeguard systems	Azure OpenAI: “Standard REST API,” “Accelerated Inference”	Identity and safety governance is fully documented , but not present in inference marketing —creates dual narrative between access control and execution.
NVIDIA	Reflex-based gating, token/chunk policy streaming, execution filters	NeMo Guardrails docs, GitHub, streaming config, Orchestrator references	TensorRT, H100/H200, NIM: “Optimized for performance,” “Agentic AI inference”	Reflex gating is buried in Guardrails layer , while accelerator docs omit governed behavior , creating semantic disaggregation .
OpenAI	Moderation, output suppression, auth tokens, scoped access	Moderation API, content policies, cookbook examples, platform documentation	API: “Simple REST interface,” “High-speed, low-cost inference,” “Enterprise-ready”	Moderation behaviors are real and reflexive , but inference is sold as non-governed , hiding session and content gating behind separate documentation.

Anthropic	Safety rail architecture, policy-based halting, permission gating	Claude safeguard policies, safety center updates, TIME article, Agency Journal	Claude API: "Drop-in compatibility," "Helpful, harmless AI," "No fine-tuning needed"	Gating and safety reflexes exist in policy layers , but inference APIs positioned as frictionless and ungated .
Intel	(No surfaced governance primitives)	Gaudi v1.22 docs, SynapseAI repo, performance sheets, Habana release decks	"Gaudi2: Best-in-class AI throughput," "MLPerf verified," "DeepSpeed optimization"	Acts as negative control . No TTL, no credential scoping, no reflex loop = proof of absence , reinforcing contrast with governed competitors.

Let me know when you've copied this, and I'll move to the **SSPS Function** → **Vendor Mapping Table** next. Just say:

"Ready for the next one."

SOUTHERN
STAR